

Operating Systems Design

Project 2 Design Document

Task I

(30%, 125 LINES) IMPLEMENT THE FILE SYSTEM CALLS (CREATE, OPEN, READ, WRITE, CLOSE, AND UNLINK, DOCUMENTED IN SYSCALL.H). YOU WILL SEE THE CODE FOR HALT IN USERPROCESS.JAVA; IT IS BEST FOR YOU TO PLACE YOUR NEW SYSTEM CALLS HERE TOO. NOTE THAT YOU ARE *NOT* IMPLEMENTING A FILE SYSTEM; RATHER, YOU ARE SIMPLY GIVING USER PROCESSES THE ABILITY TO ACCESS A FILE SYSTEM THAT WE HAVE IMPLEMENTED FOR YOU.

Implementation

First off we're going to have to implement `unloadSections()` in order to get all the rest to work. `unloadSections()` frees up the specified page out of the page table.

```
for each translation entry in the page table
    take each entry out of the page table
    if that entry is valid
        free that page
```

`handleSysCall()` takes five ints as parameters, here we're calling them `syscall`, `a0`, `a1`, `a2`, `a3`. What happens here is based on the inputs, this function figures out which further functions to direct the process to. The way that the required function is narrowed down is by comparing `syscall` to some predetermined variables. Now it's kind of hard to give pseudocode here without making it be actual code, but here we go:

```
switch (syscall)
case syscallHalt
    return handleHalt()
case syscallExit
    return handleExit(a0)
case syscallExec
    return handleExec(a0, a1, a2)
case syscallJoin
    return handleJoin(a0)
case syscallCreate
    return handleCreate(a0)
case syscallOpen
    return handleOpen(a0)
case syscallRead
    return handleRead(a0, a1, a2)
case syscallWrite
    return handleWrite(a0, a1, a2)
case syscallClose
    return handleClose(a0)
case syscallUnlink
    return handleUnlink(a0)
```

otherwise

 set an alarm saying there's an unknown syscall

 handleCreate() takes an integer (name) as a parameter, and from there a file from the VM and puts it into the storage. After putting it into storage, the filename is inserted into the fileTable.

 read the filename from virtual memory (looked up from name)

 open the filename (put it into a variable named file)

 if file isn't null

 for (2 to the length of fileTable)

 if fileTable at that spot is empty

 insert file into fileTable at that spot

 return i

 handleOpen() takes an integer (name) as a parameter. With that it does the exact same thing as handleCreate() did. As far as the filesystem is concerned they are the exact same thing.

 handleRead() takes three ints as parameters (fd, buffer, and s). It takes a file out of storage and moves it into the VM.

 grab entry from fileTable at spot fd

 if the file is null

 return -1 (we have a problem)

 byte[] tempbuff = array of bytes[s]

 int readSize = read a file(tempbuff, 0, s)

 write into VM (buffer, tempbuff)

 return readSize

 handleWrite() takes three ints as parameters (fd, buffer, s). This one takes a file from VM and moves it into storage.

 grab entry from fileTable at spot fd

 if the file is null

 return -1 (we have a problem)

 byte[] tempbuff = array of bytes[s]

 read from VM(buffer, tempbuff)

 return file.write(tempbuff, 0, s)

 handleClose() Takes a file descriptor (int). Call close() on that location in your file table and set it to null.

 grab entry from fileTable at spot fd

 if the file is not null

 call close on fileTable[fd]

 set fileTable[fd] to null

 return 0

 handleUnlink(): Takes a file descriptor (int). From here send error notifiers if the file system can not successfully unlink the file.

 try

 String name = read from VM(fd, STRINGSIZE)

 if(name is null or its length is 0)

```

        return -1 (we have a problem)
    make a new file system variable (fs) set to
Machine.stubFileSystem()
    if fs.remove(name)
        return 0 (it worked!)
    else
        return -1 (we have a problem)
catch(Exception e)
    return -1

```

Task II

(25%, 100 LINES) IMPLEMENT SUPPORT FOR MULTIPROGRAMMING. THE CODE WE HAVE GIVEN YOU IS RESTRICTED TO RUNNING ONE USER PROCESS AT A TIME; YOUR JOB IS TO MAKE IT WORK FOR MULTIPLE USER PROCESSES.

Implementation

readVirtualMemory() takes three ints (vaddr, offset, and length) and an array of bytes (data). For what needs to be added to the existing functions, we need to figure out where in the virtual memory the item is sitting, make sure that there's nothing wrong with its physical address, and copy it from the physical address.

```

int vpn = vaddr / pageSize
int voffset = vaddr % pageSize
make a translation entry (entry) from the pageTable (vpn)
set entry's used state to true
int physAddr = entry.ppn*pageSize + voffset

if (physAddr < 0 or physAddr >= memory.length or entry is invalid)
    return 0 (we have a problem)

int amount = Math.min(length, memory.length-physAddr)
System.arraycopy(memory, physAddr, data, offset, amount)

return amount

```

writeVirtualMemory() just like the read counterpart, it takes three ints (vaddr, offset, and length) and a byte array (data). For what needed to be added, most of it was the same as the read one, but instead of moving it *from* physical to, we're moving it *to* the physical memory.

```

int vpn = vaddr / pageSize
int vpoffset = vaddr % pageSize
make a translation entry (entry) from the pageTable (vpn)
set the entry's used state to true
calculate the physical address passed on the pageSize + voffset

if (physAddr < 0 or physAddr >= memory.length or entry is invalid)
    return 0 (we have a problem)

```

make sure the entry's dirty bit is set to true

```

make an int (amount) set to the min of length and memory.length-
physAddr
System.arraycopy(data, offset, memory, physaddr, amount)

```

```

    loadSections() for what we needed to implement, within the for loop (within the other for-
loop) we had to take the given address and follow it to the actual address.
make an int (vpn) set to section.getFirstVPN()+1
make a new translation entry (entry) set to pageTable[vpn]
entry.ppn = UserKernel.availablePages.removeFirst()
set entry.valid to true
entry.readOnly = section.isReadOnly()
section.loadPage(loop variable i, entry.ppn)

```

Task III

(30%, 125 LINES) IMPLEMENT THE SYSTEM CALLS (EXEC, JOIN, AND EXIT, ALSO DOCUMENTED IN SYSCALL.H).

Implementation

handleExec takes three parameters, all ints (name, argc, argv). Then it takes instructions out of VM and lines them up for execution in a child thread.

```

make an array of strings (args) set to newString[argc]
for i < argc
    byte[] pointer = new byte[ADDRESSSIZE]
    readVirtualMemory(argv + i*ADDRESSSIZE, pointer)
    args[i] = readVirtualMemoryString(Lib.bytesToInt(pointer, 0),
STRINGSIZE)
make a new userProcess(child)
add child to myChildren
make a string(childname) set to readVirtualMemoryString(name,
STRINGSIZE)
return child.execute (childname, args) ? child.myPID : -1

```

handleJoin() takes one int parameter (pid). It searches through all the children, and if that child's pID matches the one that needs to be joined, it is.

```

for every child in myChildren
    if the child's pID matches the one we're looking for
        mychildren.remove(child)
        return child.myExitStatus
return -1

```

handleExit takes one int parameter, (exit). It tries to shut down all the threads, halts the machine if it can, calls KThread.finish(), and returns the exit number.

```

set myExitStatus to exit
call unloadSections()
for every open file in the file table
    if the file is not null
        call file.close()
if myPID equals 0

```

```

        call machine.halt()
    KThread.finish()
    return exit

```

Task IV

(15%, 50 LINES+EXISTING PRIORITY SCHEDULER) IMPLEMENT A LOTTERY SCHEDULER (PLACE IT IN THREADS/LOTTERYSCHEDULER.JAVA). NOTE THAT THIS CLASS EXTENDS PRIORITYSCHEDULER, YOU SHOULD BE ABLE TO REUSE MOST OF THE FUNCTIONALITY OF THAT CLASS; THE LOTTERY SCHEDULER SHOULD NOT BE A LARGE AMOUNT OF ADDITIONAL CODE. THE ONLY MAJOR DIFFERENCE IS THE MECHANISM USED TO PICK A THREAD FROM A QUEUE: A LOTTERY IS HELD, INSTEAD OF JUST PICKING THE THREAD WITH THE MOST PRIORITY. YOUR LOTTERY SCHEDULER SHOULD IMPLEMENT PRIORITY DONATION. (NOTE THAT SINCE THIS IS A LOTTERY SCHEDULER, PRIORITY INVERSION CAN'T ACTUALLY LEAD TO STARVATION! HOWEVER, YOUR SCHEDULER MUST DO PRIORITY DONATION ANYWAY.)

Implementation

In implementing the lottery scheduler, the necessary methods added were `setPriority`, `increasePriority()`, `decreasePriority()`, and `getEffectivePriority()`. There is also the constructor `LotteryScheduler()` which doesn't do anything.

`setPriority()` takes two parameters, the first is a `KThread` (`thread`) and an `int` (`priority`). It then makes sure that interrupts are disabled, and sets the priority.

make sure that interrupts are disabled

make sure that the priority is greater than or equal to the `MINPRIORITY`

set the thread state to priority

`increasePriority()` takes no parameters. First it disables interrupts, makes sure that the priority isn't the max priority (if it is it returns false) and sets the priority to one higher, and re-enables interrupts (then returns true).

disable interrupts

grab the current thread (`thread`)

grab thread's priority (`priority`)

if `priority` is the `MAXPRIORITY`

return false

set thread's priority to `priority + 1`

re-enable interrupts

return true

`decreasePriority()` takes no parameters. It then basically does the exact same thing that `increase` did, but in the opposite direction.

disable interrupts

grab the current thread (`thread`)

grab thread's priority (`priority`)

if `priority` is the `MINPRIORITY`

return false

set thread's priority to `priority - 1`

re-enable interrupts

```
return true
```

getEffectivePriority() takes one parameter, a ThreadState (mystate). It then grabs mystates's priority, goes through the waiting queue inside the donation queue and if the current thread's state is the same as mystate, we continue, otherwise adds the effective priority to mystate's, and ends up returning it.

```
grab the priority of mystate and put it into variable tickets
for every priority queue (myQueue) in mysteries' donation queue
    for every KThread (currentThread) in myQueue's waiting queue
        if currentThread's thread state is equal to mistake
            continue
        tickets += currentThread's effective priority
return tickets
```